

04-630
Data Structures and Algorithms for Engineers

George Okeyo
Carnegie Mellon University Africa
gokeyo@andrew.cmu.edu

Lecture 17: Graphs

Outline

- Graphs: Preliminaries
 - Applications
 - Definitions
- Graph representation:
 - adjacency matrix
 - adjacency list
- Graph traversal
 - BFS, DFS
- Applications of graph search algorithms


Graphs: Applications

- Mapping
- Transportation
- Electrical engineering
- Computer networks

Graphs: Definitions

- A graph **G** , formally **$G=(V,E)$** , is composed of a set of ***vertices*** **V** and a set of ***edges*** **$E \subset V \times V$** .
 - A ***vertex*** refers to a node in a graph.
 - ***Edges*** connect vertices.
- An edge **$e=(u,v)$** is a pair of vertices for ***directed*** graphs.
- For an undirected graph, an edge between **\underline{u}** and **\underline{v}** is represented by two pairs **$(u,v) \in E$** and **$(v,u) \in E$** .

Graphs: directed graphs (digraphs)



$$\mathbf{V} = \{A, B, C, D, E\}$$


$$\mathbf{E} = \{(A, B), (A, D), (B, D), (D, C), (D, E)\}$$

Has ***directed edges***.

An edge (u, v) is **directed** if the pair (u, v) is ***ordered*** such that u precedes v .

End points: u and v . In the case of ***directed graphs***, u is ***origin*** and v is the ***destination***.

Graphs: undirected graphs



Has *undirected edges*.

An edge (u,v) is undirected if the pair (u,v) is *not ordered* such that u precedes v .


$V = \{A, B, C, D, E\}$

$E = \{(A,B), (B,A), (A,D), (D,A), (B,D), (D,B), (D,C), (C,D), (D,E), (E,D)\}$

Or


$E = \{\{A,B\}, \{A,D\}, \{B,D\}, \{D,C\}, \{D,E\}\}$, set of all 2-element subsets of V

1. It is also possible to have a *mixed graph*, with both directed and undirected edges.
2. Both undirected and mixed graphs can be converted (*if necessary*) to a directed graph.



An undirected graph converted to a directed graph.

Graphs: mixed graph example



Exercise: 1) Roads in Kigali? What kind of scenario(s) in city road network design are represented by the graph?
2) Give another example of a real-world problem that can be modelled with a mixed graph.

Graphs: definitions




- **Adjacent vertices:** Vertices u and v are adjacent if there is an edge with u and v as vertices, i.e. $(u,v) \in E$.
- **Degree** of a **vertex**: number of adjacent vertices.
- An edge e is **incident** on a vertex if the vertex is one of the edges endpoints.
- **Path**: sequence of vertices v_1, v_2, \dots, v_n , such that v_{k+1} is adjacent to v_k for $k=1 \dots n-1$.

Graphs: definitions

Simple path: a path with no repeated vertices, e.g. A, B, D.

Cycle: a simple path where the first and the last vertices are the same, e.g. A, B, D, A.

Connected graph: any two vertices are connected by some path.




Subgraph: a subset of vertices and edges forming a graph, i.e., $G'=(V',E')$ is a subgraph of $G(V,E)$ if $V' \subseteq V$, $E' \subseteq E$ and G' is a graph.

Connected component: A connected component of an undirected graph is a maximal set of vertices such that **there is a path** between **every pair of vertices**.


Tree: connected graph without cycles.

Graphs: definitions

Complete graph: every vertex is connected to every other vertex.




Complete undirected graph



Complete directed graph

Weighted graph: a graph in which every edge has an associated value or weight.




Graph representation:

A graph $G=(V,E)$ with **n vertices** and **m edges** can be represented using two data structures:

- **Adjacency matrix:** an $n \times n$ matrix M (an array), where each element
 - $M[i,j]=1$ if (i,j) is an edge of G , and
 - $M[i,j]=0$ if it is not.
 - For a weighted graph $M[i,i]=w$, the weight of the edge if (i,j) is an edge and $M[i,i]=\infty$ if it is not.
- **Adjacency list:** uses a **linked list** that stores the sequence of vertices that are adjacent to each vertex.

Graph representation: adjacency matrix

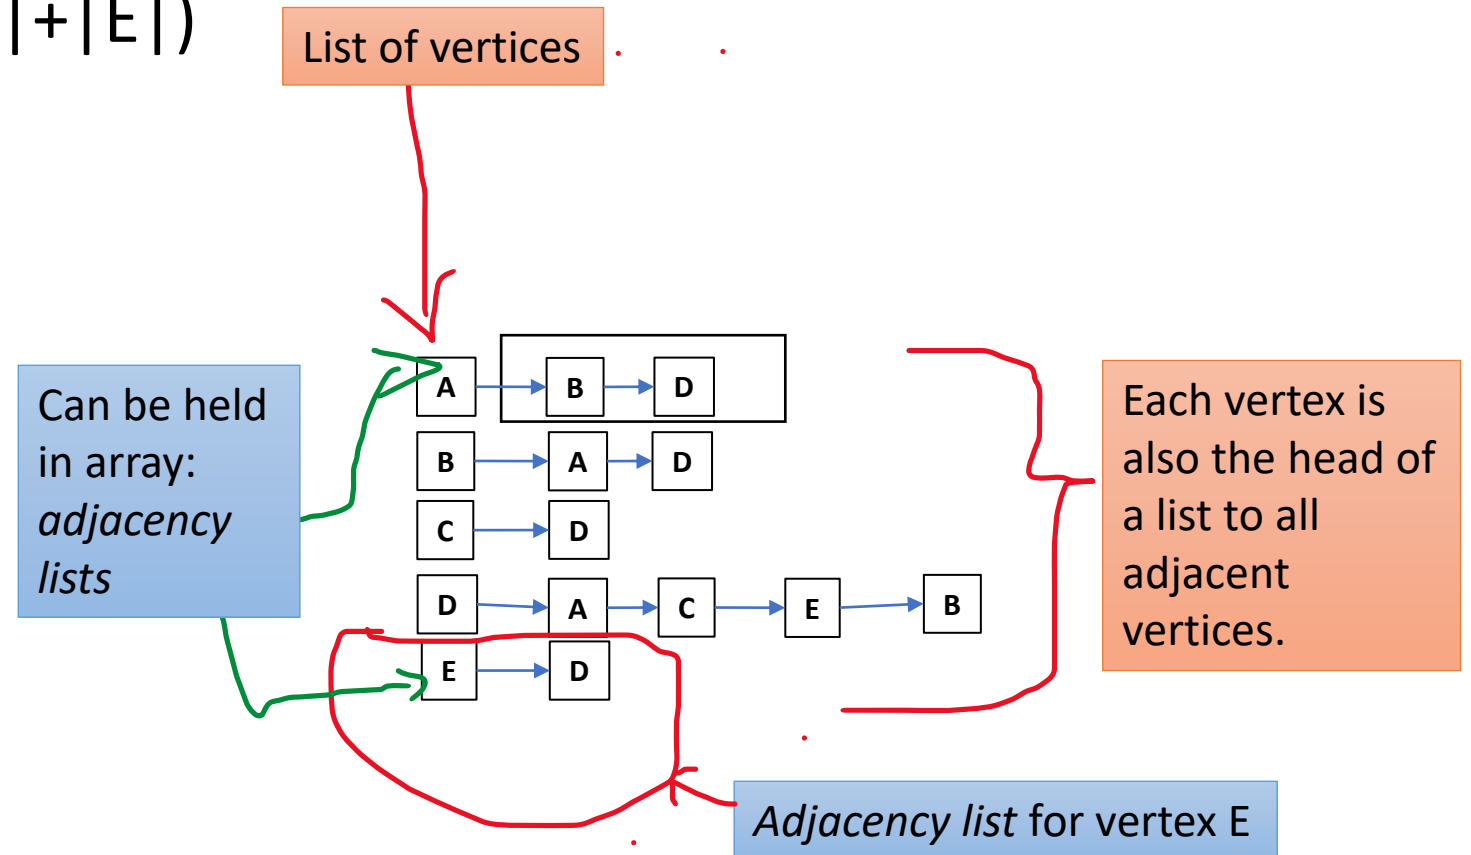

- Space requirement: $\Theta(|V|^2)$



	A	B	C	D	E
A	0	1	0	1	0
B	1	0	0	1	0
C	0	0	0	1	0
D	1	1	1	0	1
E	0	0	0	1	0

Graph representation: adjacency lists

- Space requirement: $\Theta(|V| + |E|)$



Code Implementation: Adjacency List

```
struct node
{
    int dest;
    node *next;//next node in list
};

struct adjList
{
    node *head;
};
```

```
class graph
{
private:
    int nvertices;
    adjList *adjLists;//adjacency info: list of edges
public:
    graph(int nvertices);//constructor
    node *createListNode(int dest);
    void addEdge(int srcData,int destData);
    void printGraph();
};

graph::graph(int nvertices)
{
    this->nvertices=nvertices;//set up the number of vertices
    adjLists=new adjList[nvertices];//create an adjacency list for each vertex
    for(int i=0;i<nvertices;i++)
        adjLists[i].head=NULL;//set each head of the list to NULL.
}
```

Diagram illustrating the adjacency list structure for a graph with 5 vertices:

0	1	2	3	4
head=NULL	head=NULL	head=NULL	head=NULL	head=NULL

The diagram shows a horizontal array of five boxes, each representing an adjacency list for a vertex. Each box is divided into two parts: a top part labeled 'head=NULL' and a bottom part containing the vertex index (0, 1, 2, 3, 4). An orange box labeled 'adjacency list-3' is positioned above the array, with a blue arrow pointing from it to the 'head' field of the box corresponding to vertex 3.

Code Implementation: Adjacency List

```
/*  
Utility method that creates nodes  
*/  
node *graph::createListNode(int dest)  
{  
    node * newNode=new node;  
    newNode->dest=dest;  
    newNode->next=NULL;  
    return newNode;  
}
```

Code Implementation: Adjacency List

```
/*This should create two nodes - the source and destination nodes.*/  
void graph::addEdge(int src, int dest)  
{  
    //Adding a new node  
    node * newNode=createListNode(dest);  
    newNode->next=adjLists[src].head;//make the current head of the list to be the next node for the new node  
    adjLists[src].head=newNode;//make the new node the head of the adjacency list.  
  
    //  
    newNode=createListNode(src);  
    newNode->next=adjLists[dest].head;  
    adjLists[dest].head=newNode;  
}
```


Code Implementation: Adjacency List

```
void graph::printGraph()
{
    int v;
    cout<<"Adjacency lists\n";
    for(v=0;v<nvertices;v++)//step through each adjacency list
    {
        node * aNode=adjLists[v].head;
        cout<<"\n["<<v<<"]";
        while(aNode)
        {
            cout<<"->"<<aNode->dest;
            aNode=aNode->next;
        }
        cout<<endl;
    }
}
```

Graph Traversal

- **Traversal**: systematic procedure for exploring a graph through examining all its vertices and edges.
 - Maintain information about the state of a vertex to keep track of the traversal.
 - States:
 - **Undiscovered**: the vertex is in the initial untouched state.,
 - **Discovered**: vertex has been found but its edges have not been processed.
 - **Processed**: the state of the vertex after we have visited all its edges.
- Traversal algorithms:
 - Breadth-first search (BFS)
 - Depth-first search (DFS)

Graph traversal

- Begins with a starting vertex, $s \in V$
- Coloring vertices:
 - Undiscovered vertices are colored **white**.
 - Color a vertex **gray** if it has been discovered but all its edges have not been explored.
 - Color a vertex **black** after all its adjacent vertices have been discovered.

Breadth first search (BFS)

- Uses a FIFO queue to ensure oldest unexplored vertices are processed first.
- Process:
 - The starting vertex s is assigned the distance 0.
 - All vertices that are 1 edge away from s are visited and assigned a distance 1.
 - Search proceeds to vertices 2 edges away, which are assigned distance 2.
 - Process continues until each vertex has been assigned a level/distance.
- The level of each vertex v corresponds to the length of the shortest path from s to v .

BFS algorithm

BFS(G,s):

##Initialize all vertices

For $u \in G.V$ **Do**

$u.color = white$

$u.dist = \infty$

$u.pred = NULL$

##Initialize BFS

$s.color = gray$

$s.dist = 0$

$Q = \{s\}$ ##a FIFO queue

handle all u's children

While Q is not empty **Do**

$dequeue(Q, u)$

For $v \in u.adj$ **Do**

If $v.color = white$ **Then**

$v.color = gray$


$v.dist = u.dist + 1$

$v.pred = u$

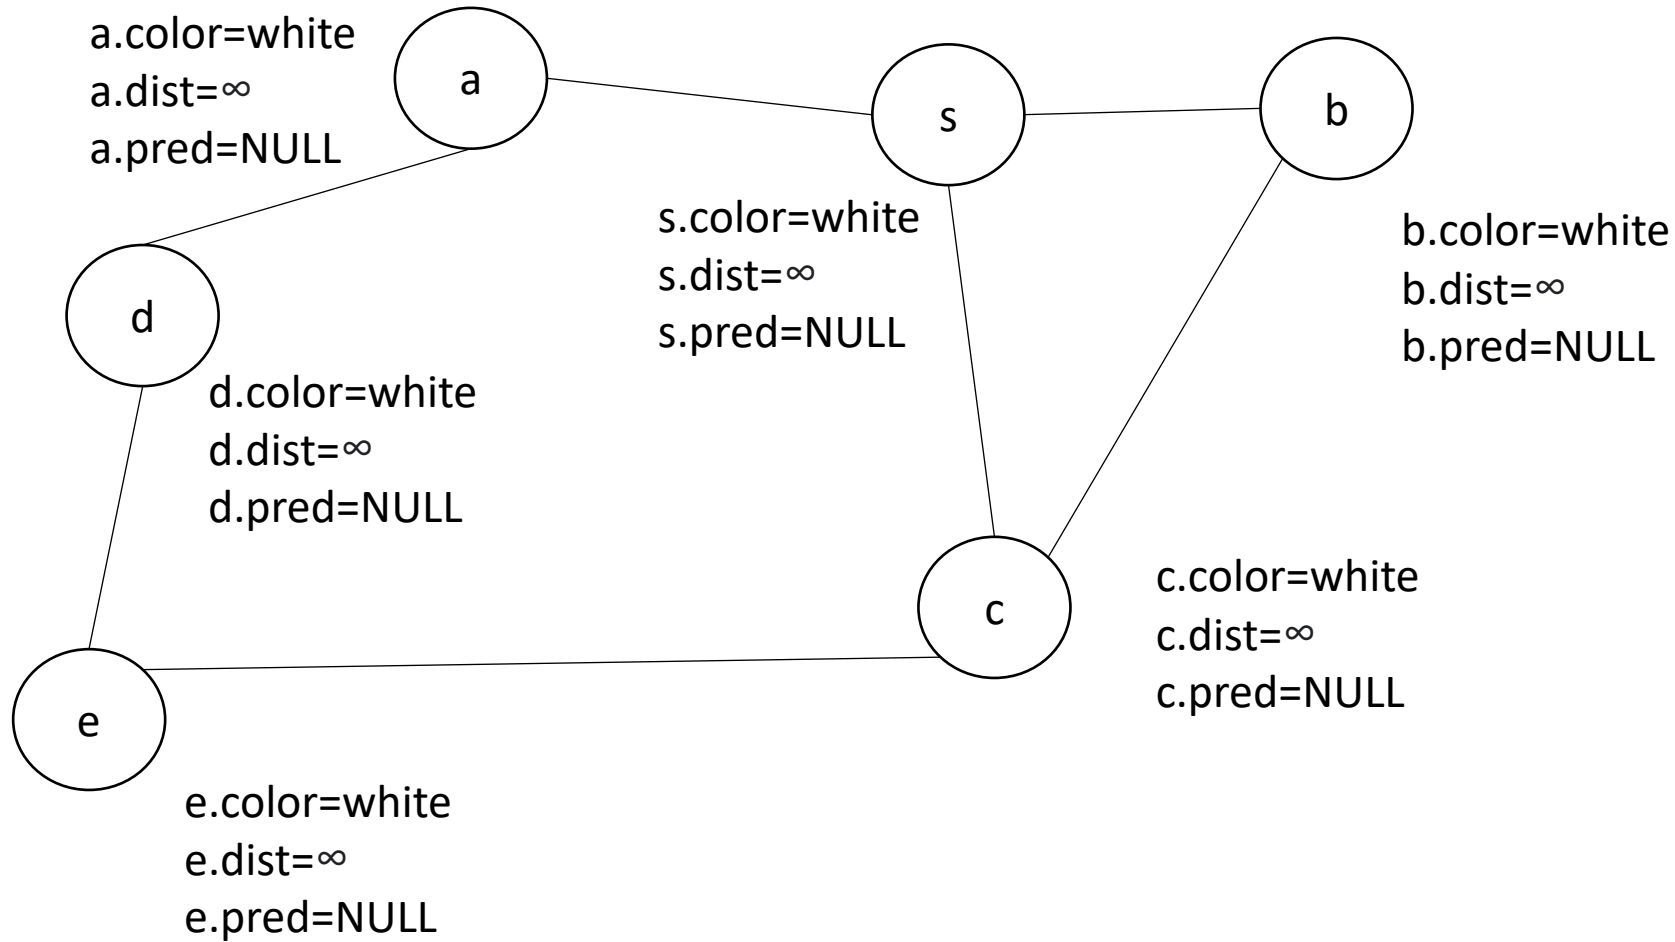
$enqueue(Q, v)$

$u.color = black$

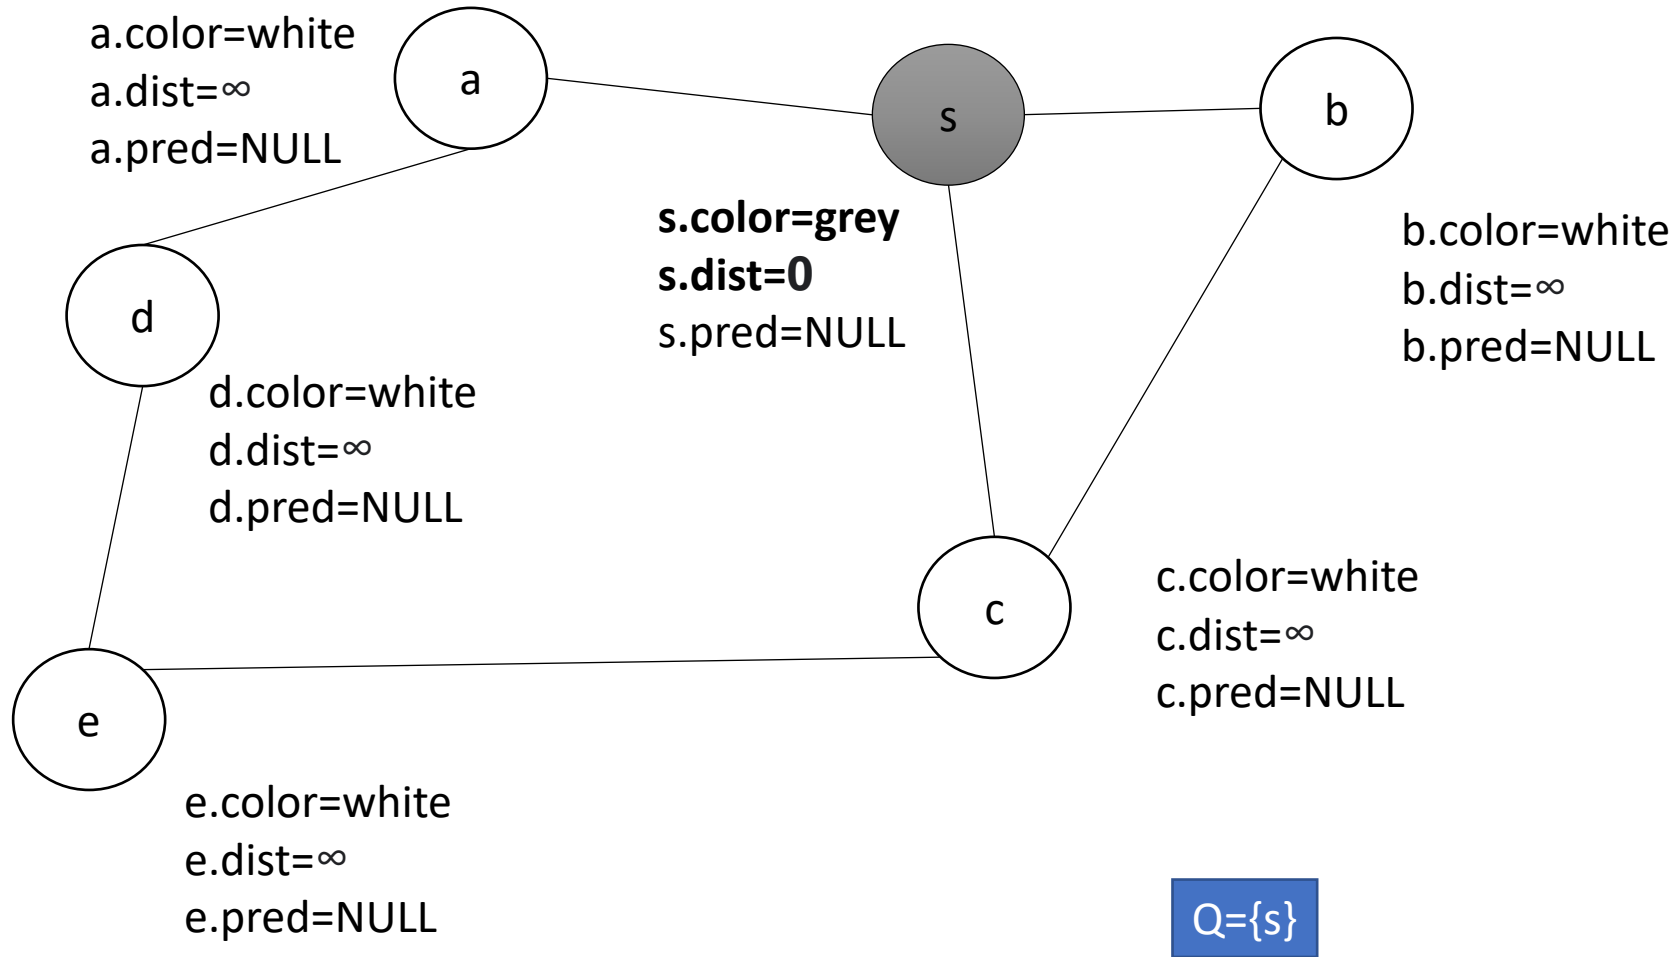
BFS Example



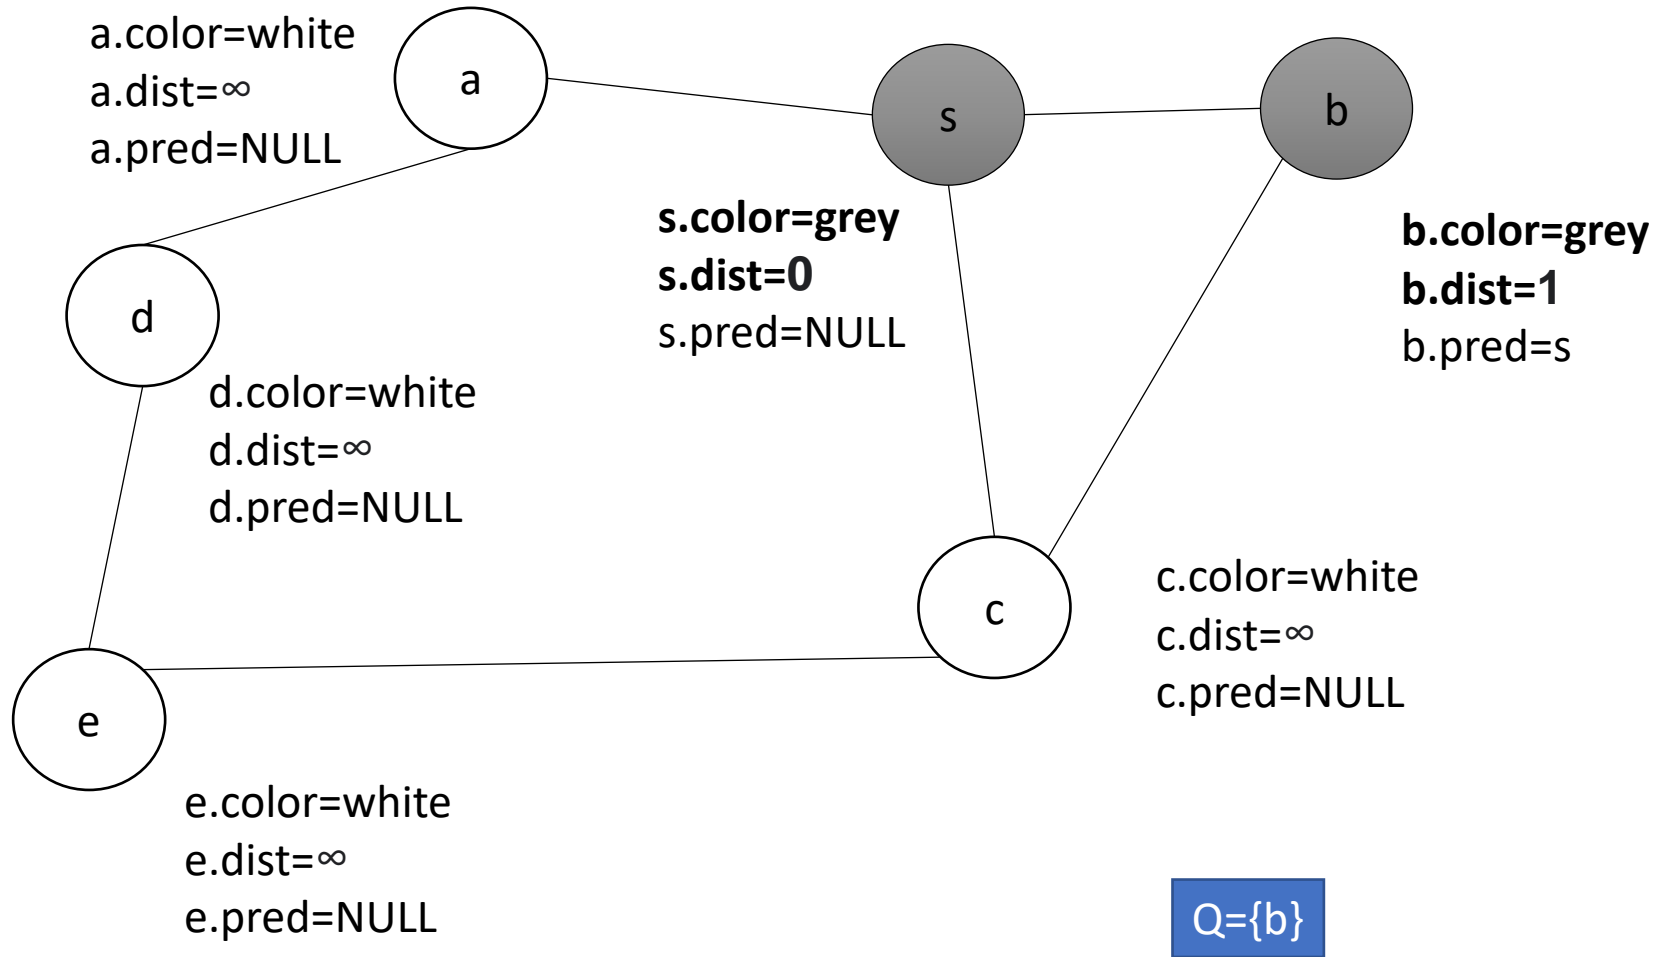
BFS Example



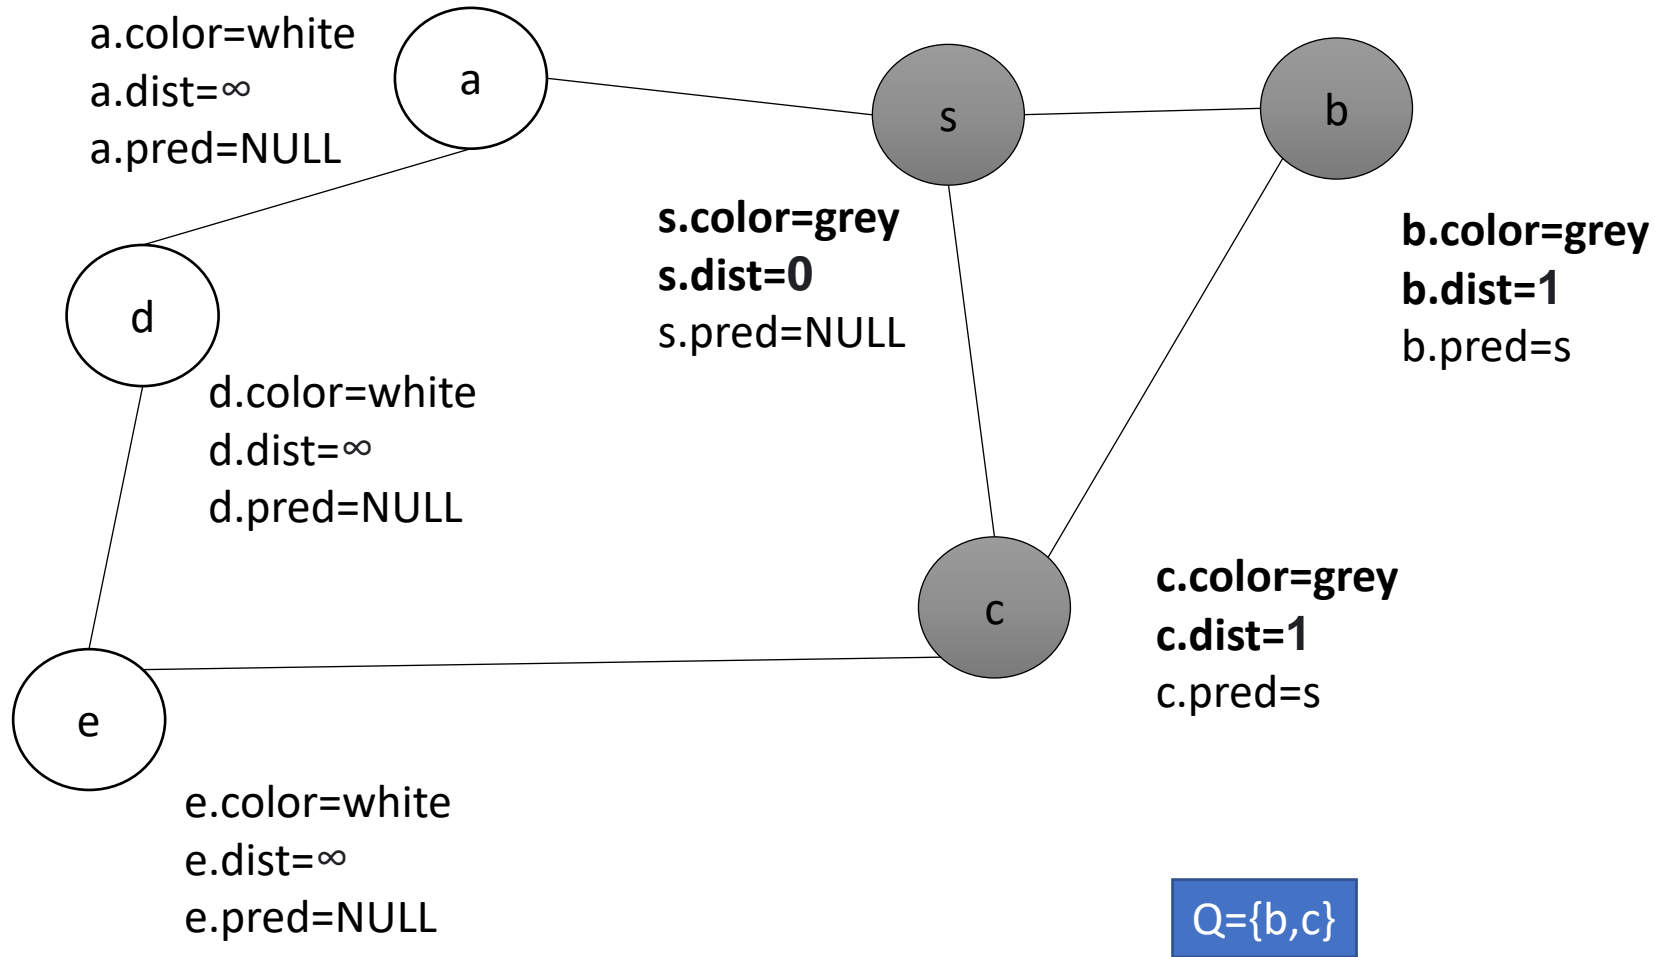
BFS Example



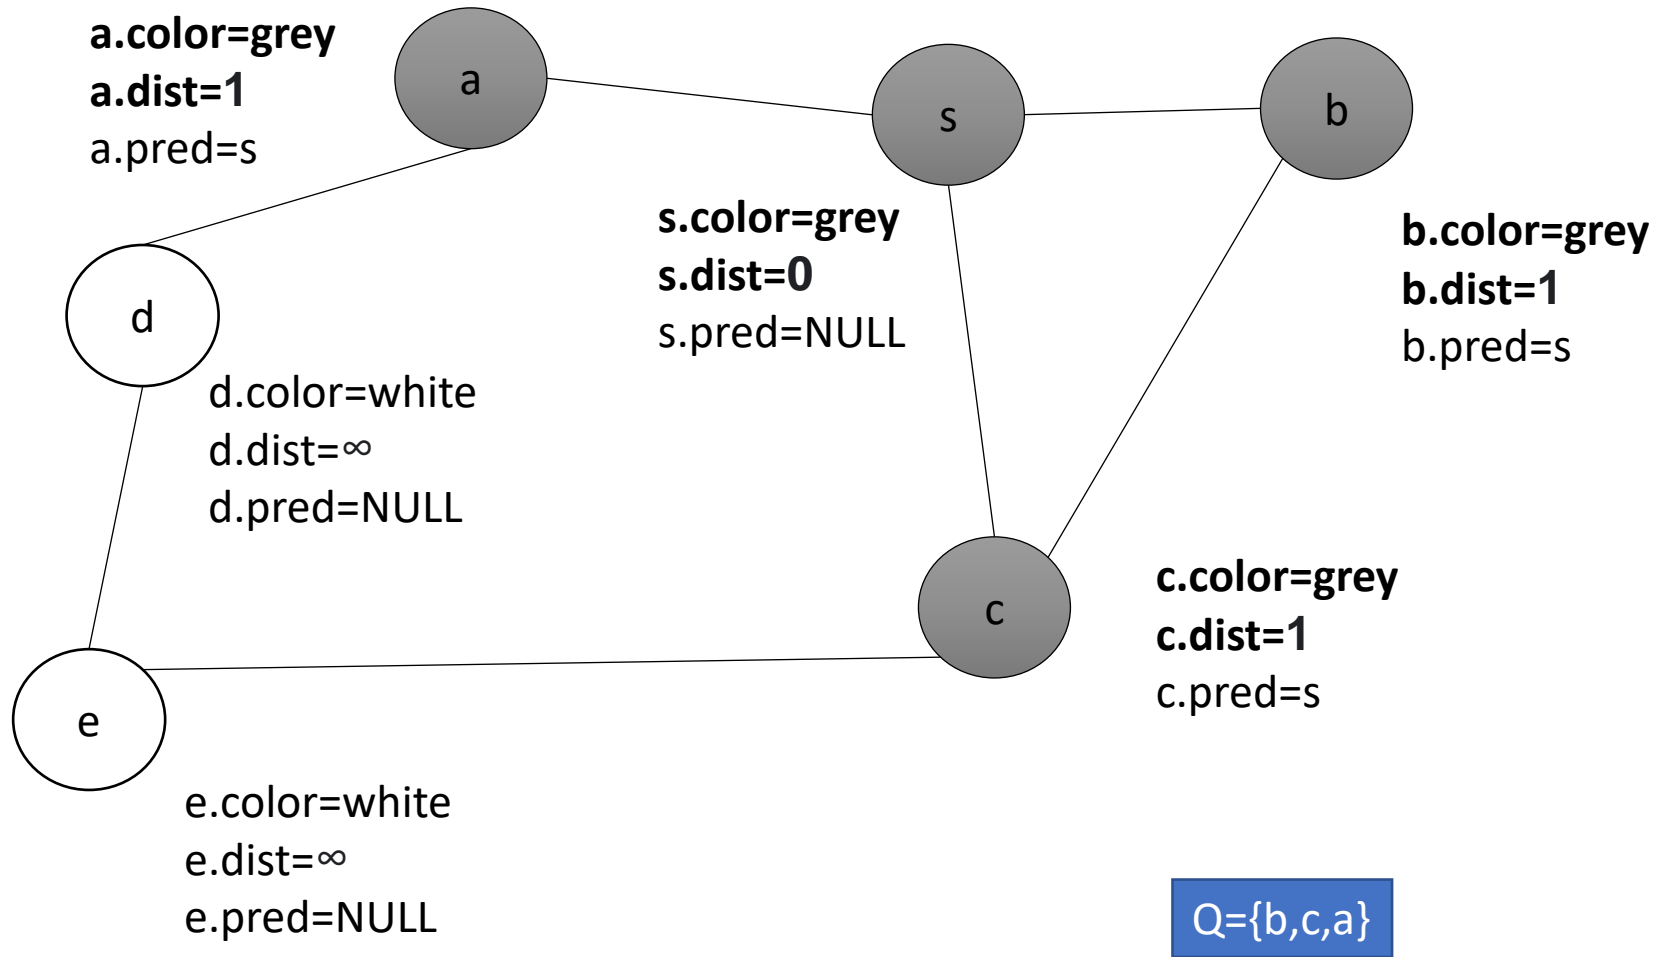
BFS Example



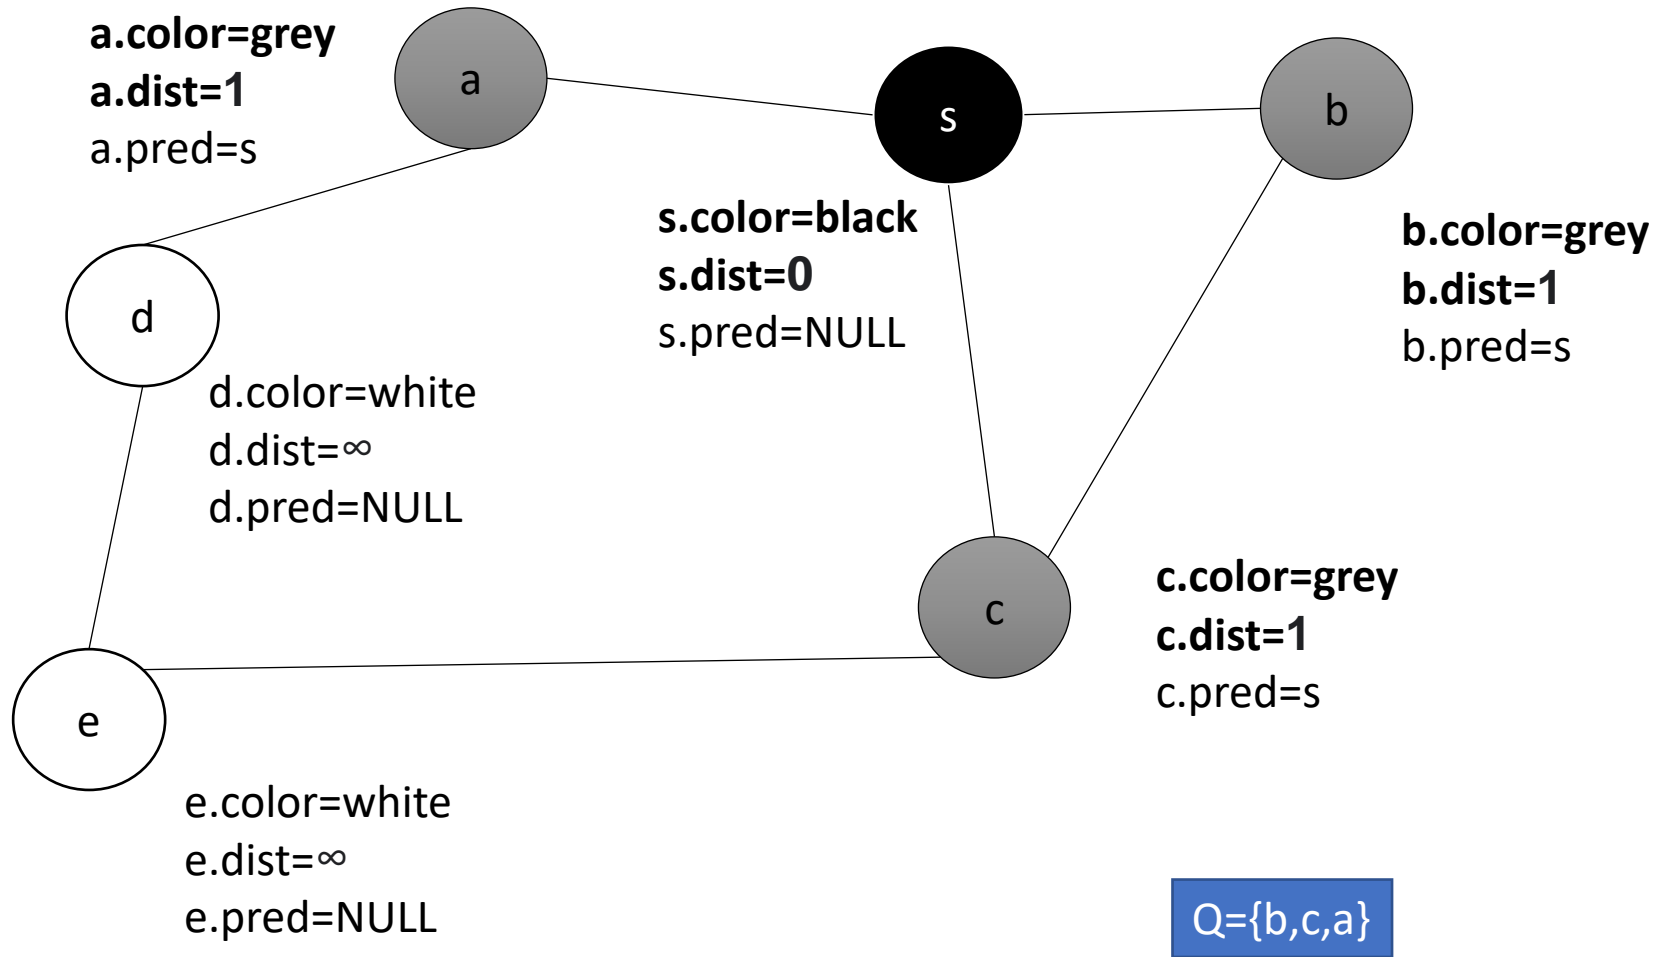
BFS Example



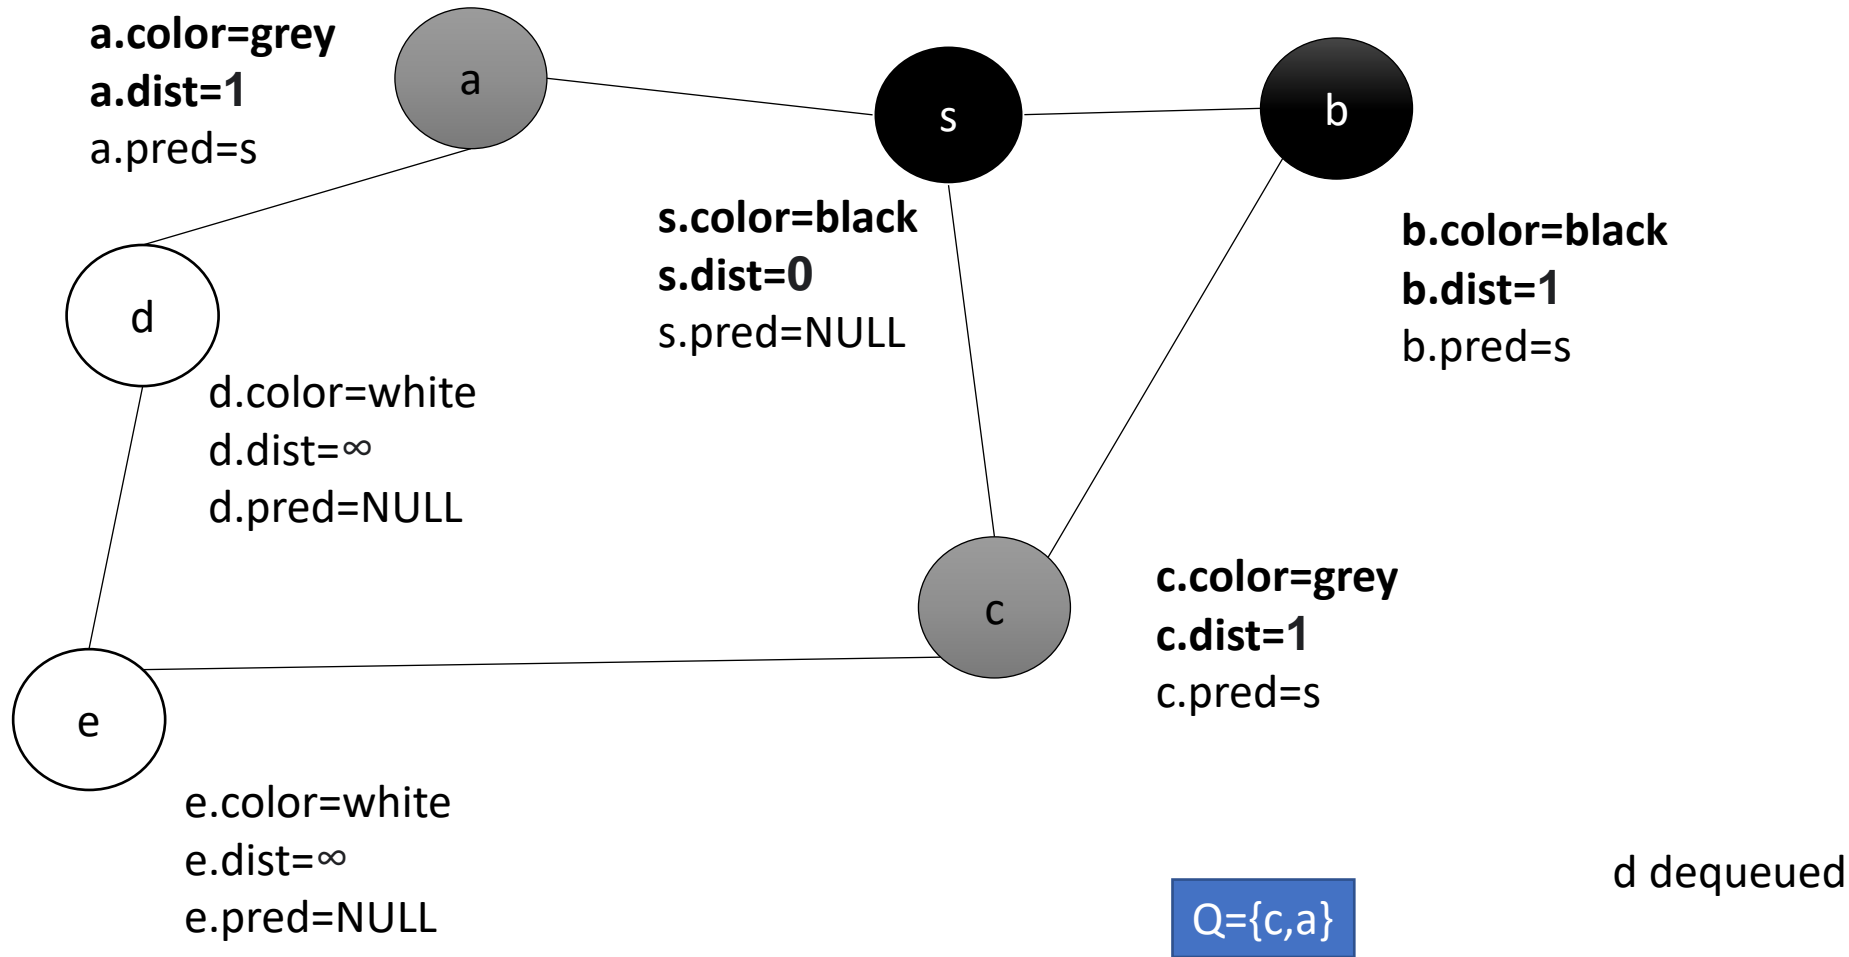
BFS Example



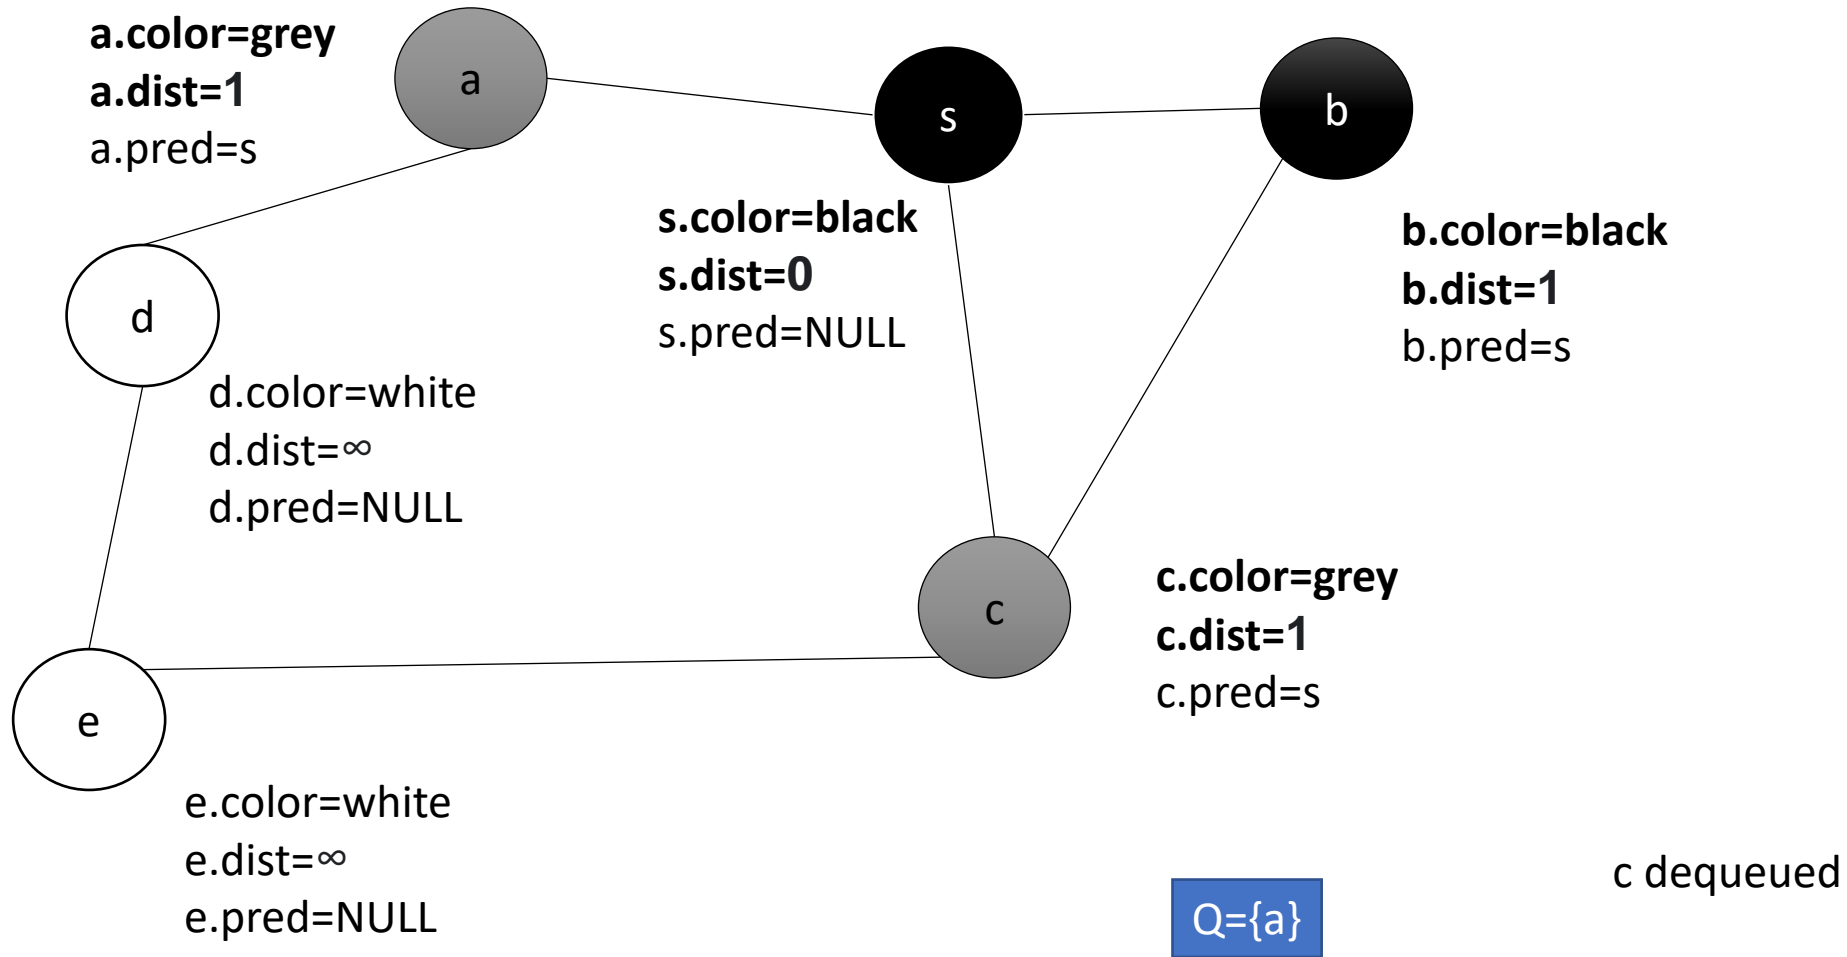
BFS Example



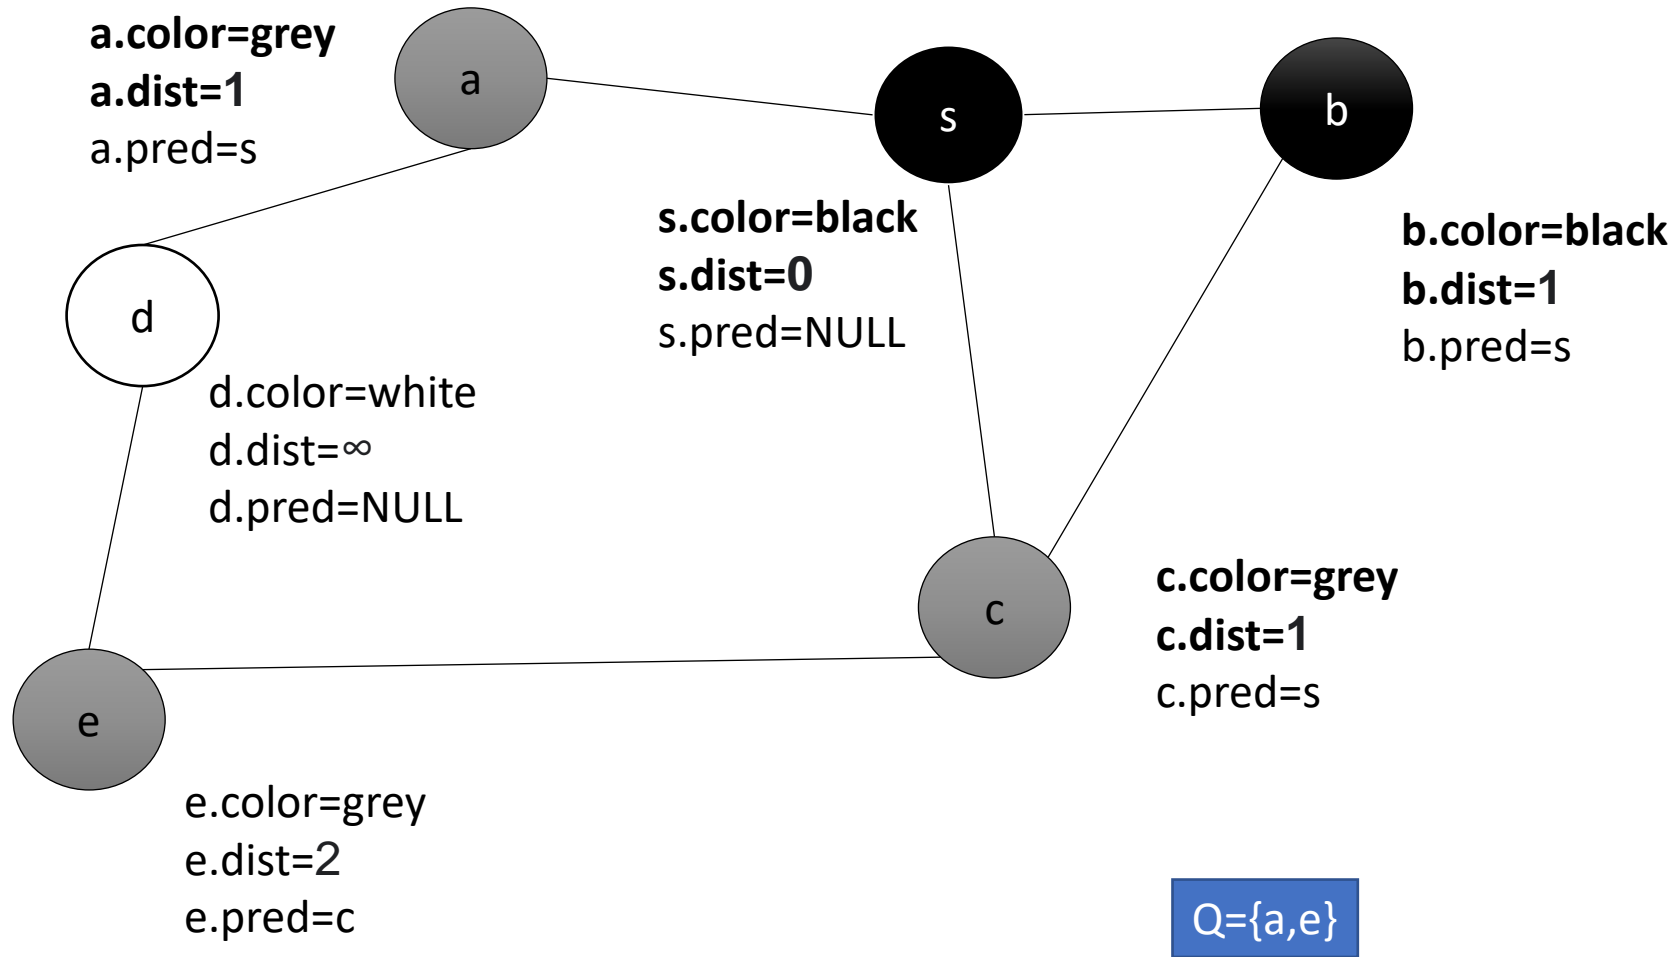
BFS Example



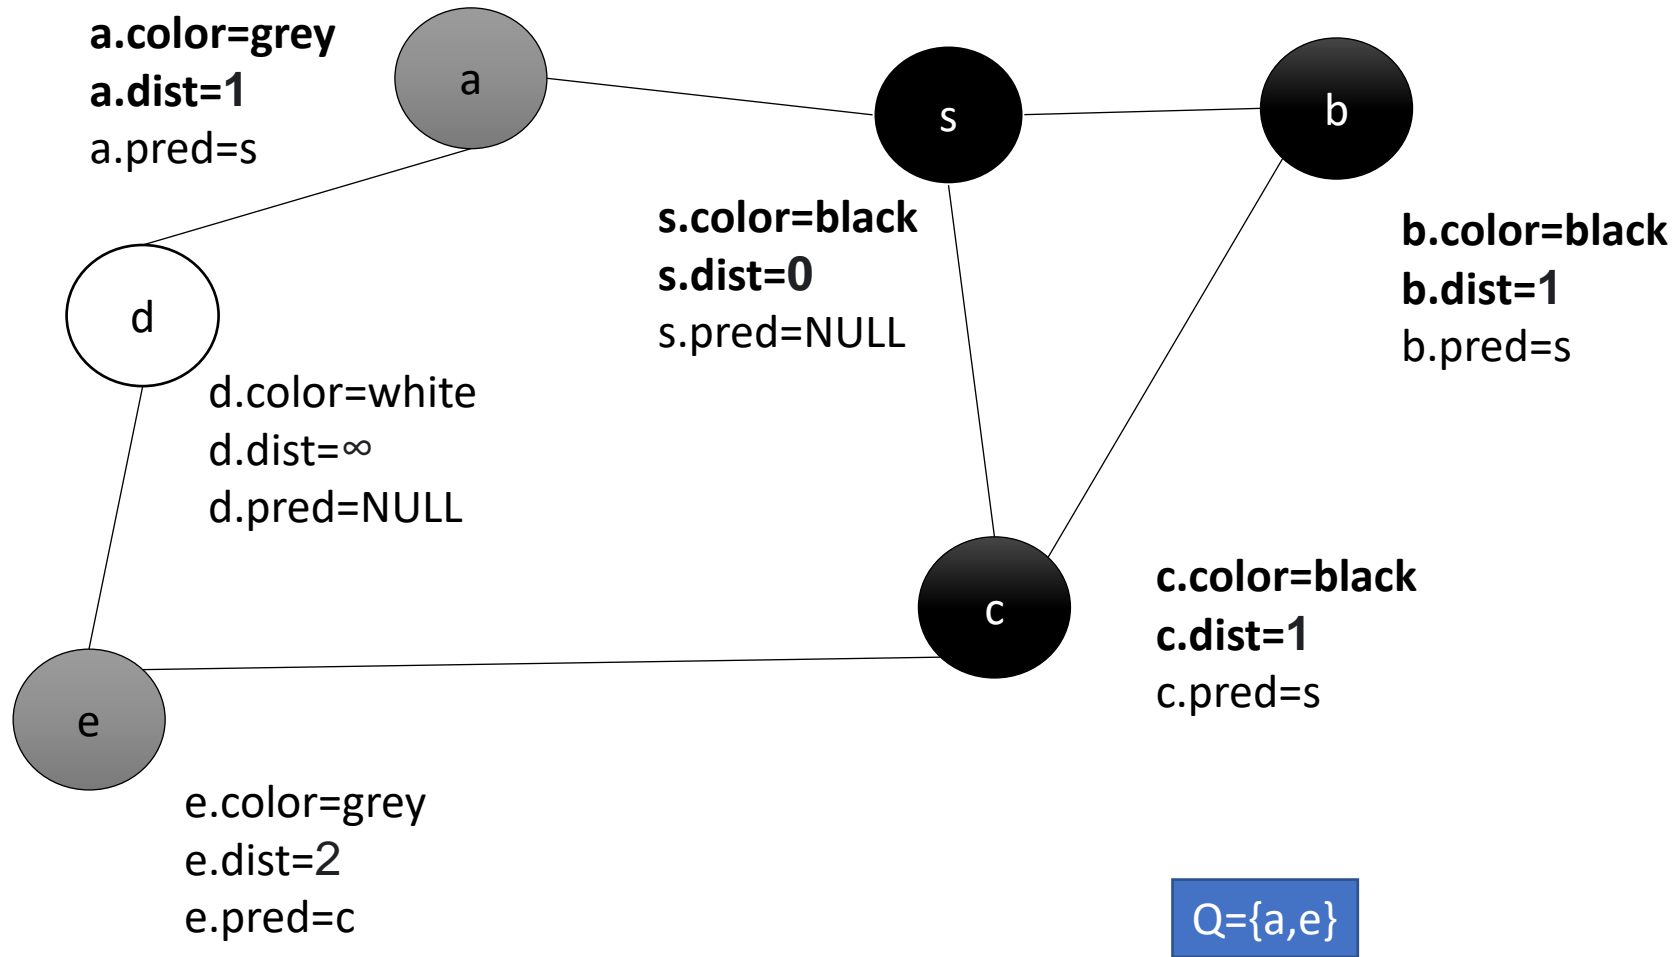
BFS Example



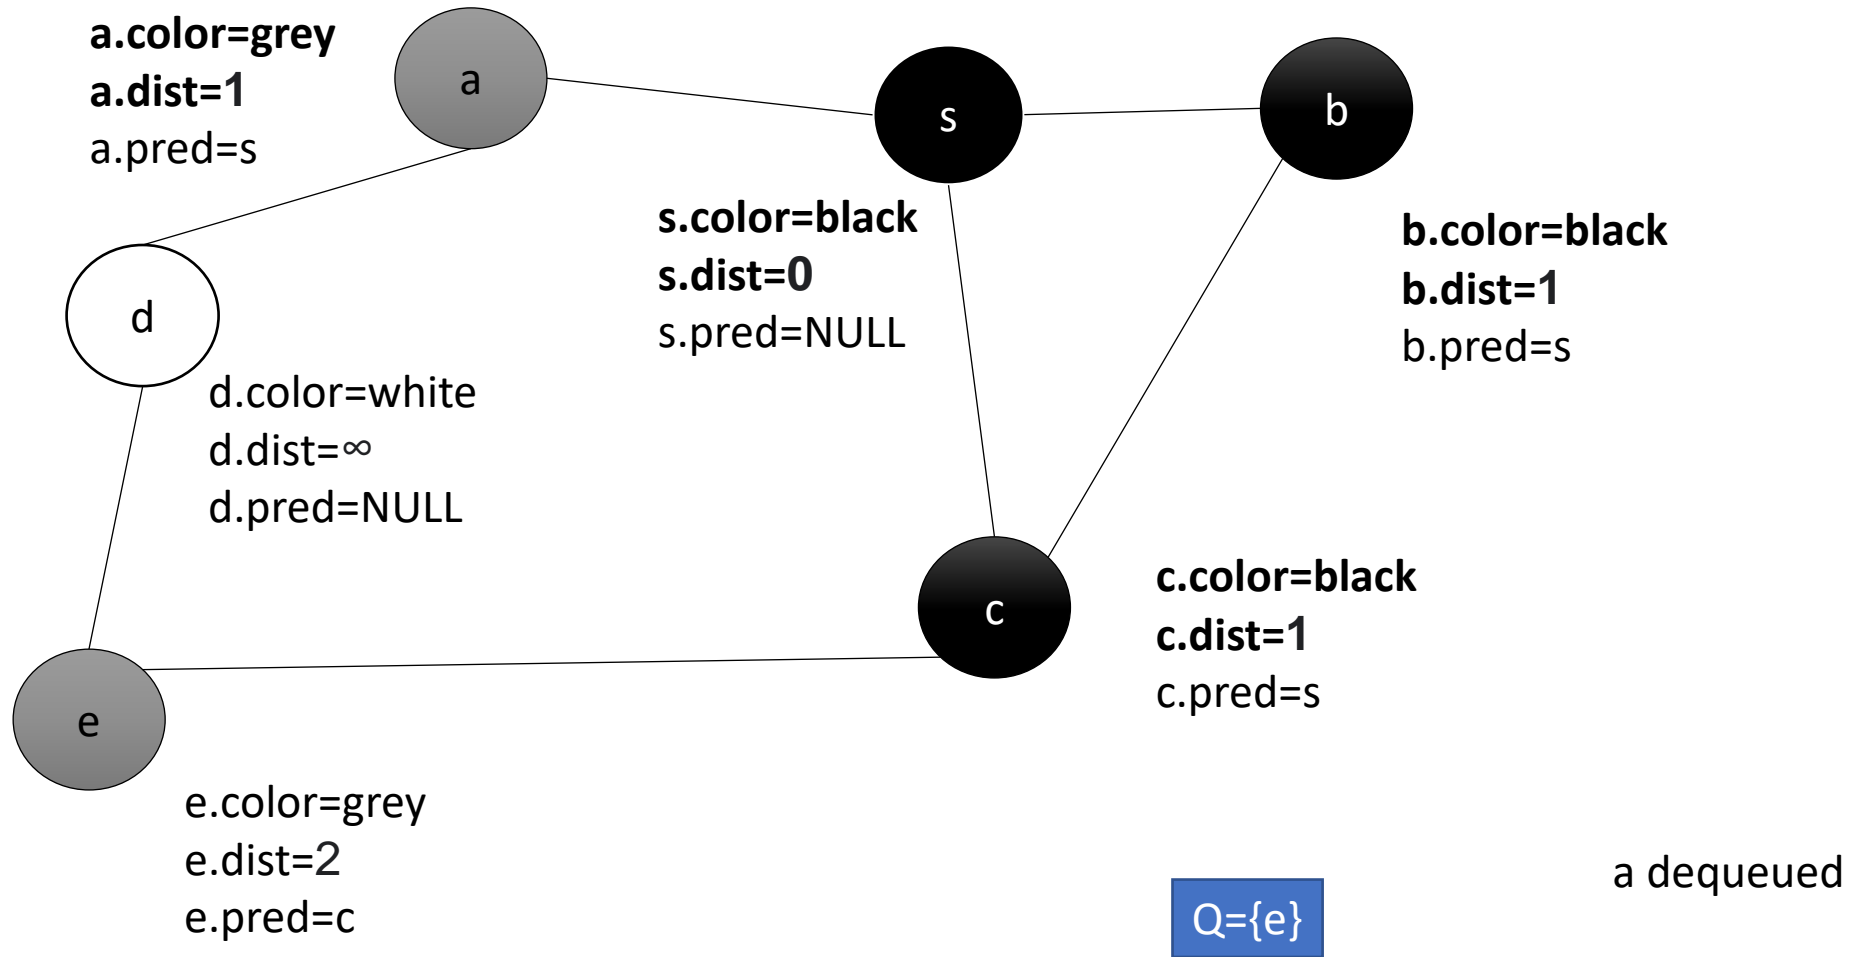
BFS Example



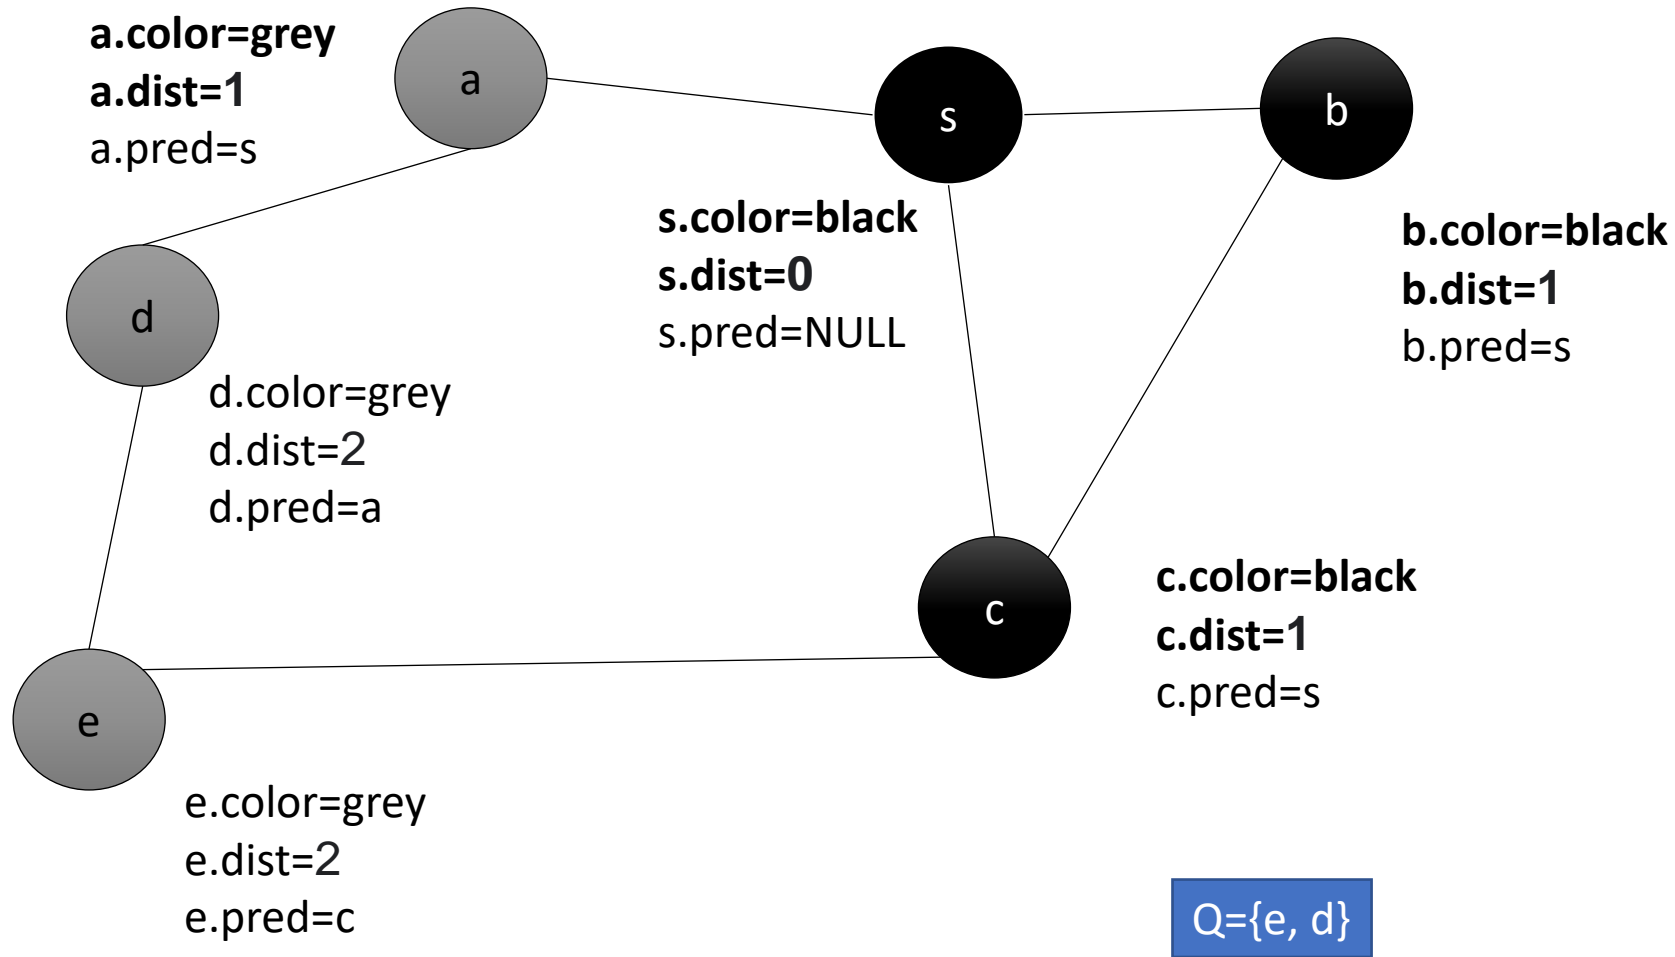
BFS Example



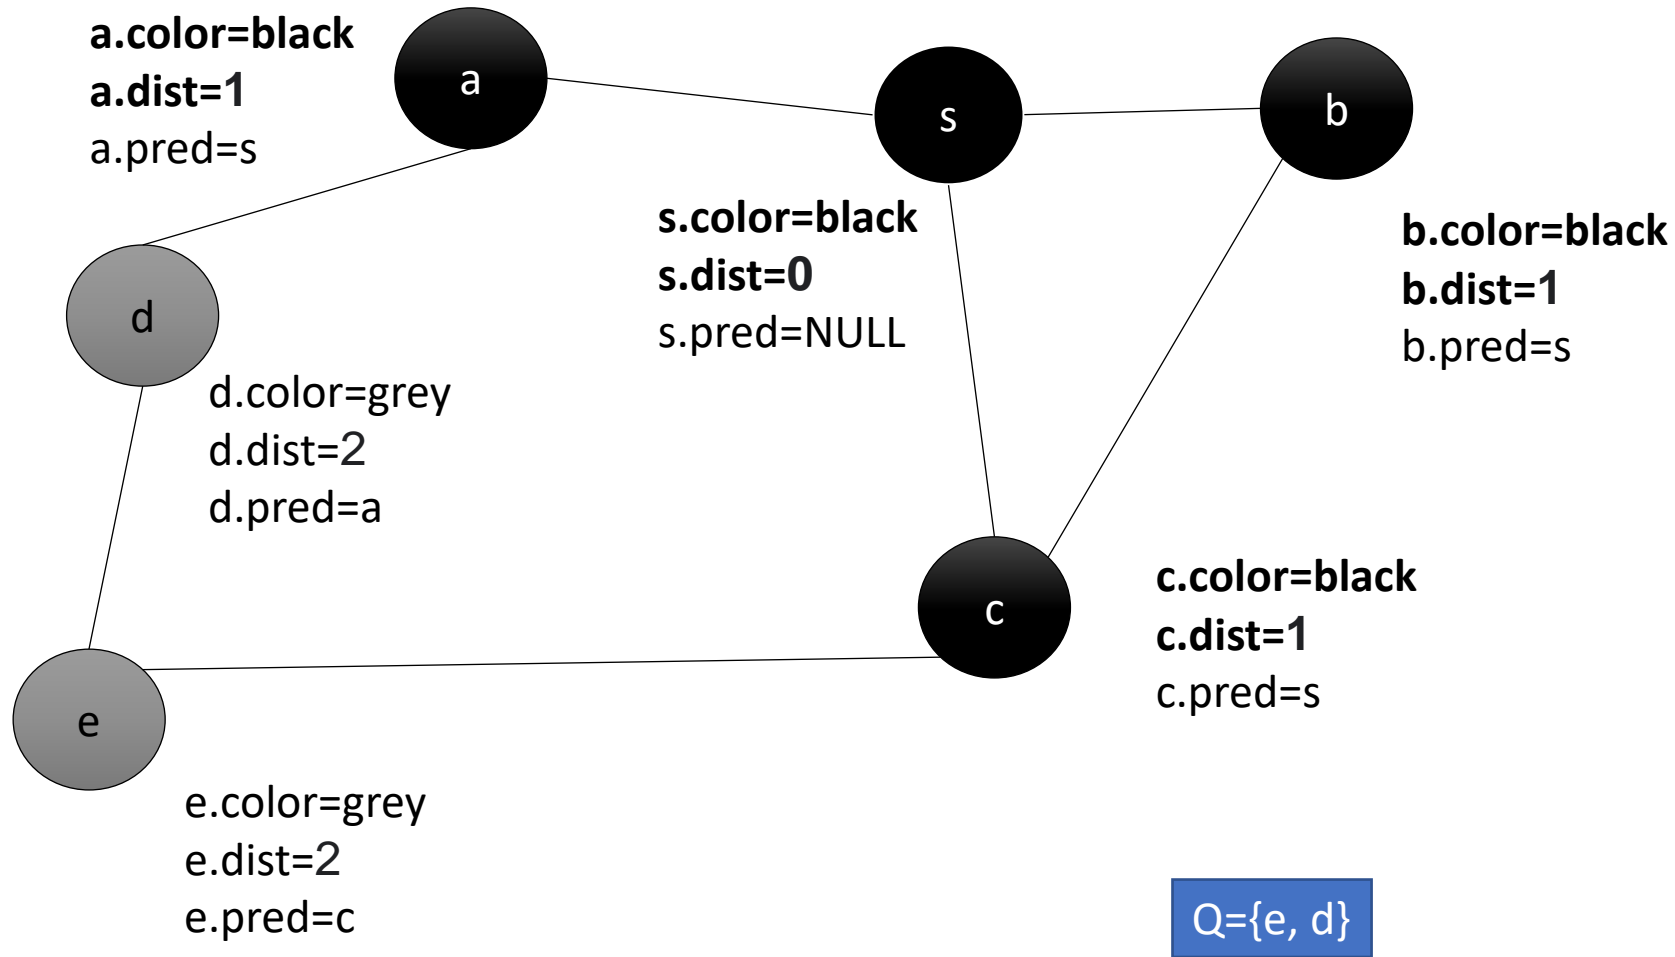
BFS Example



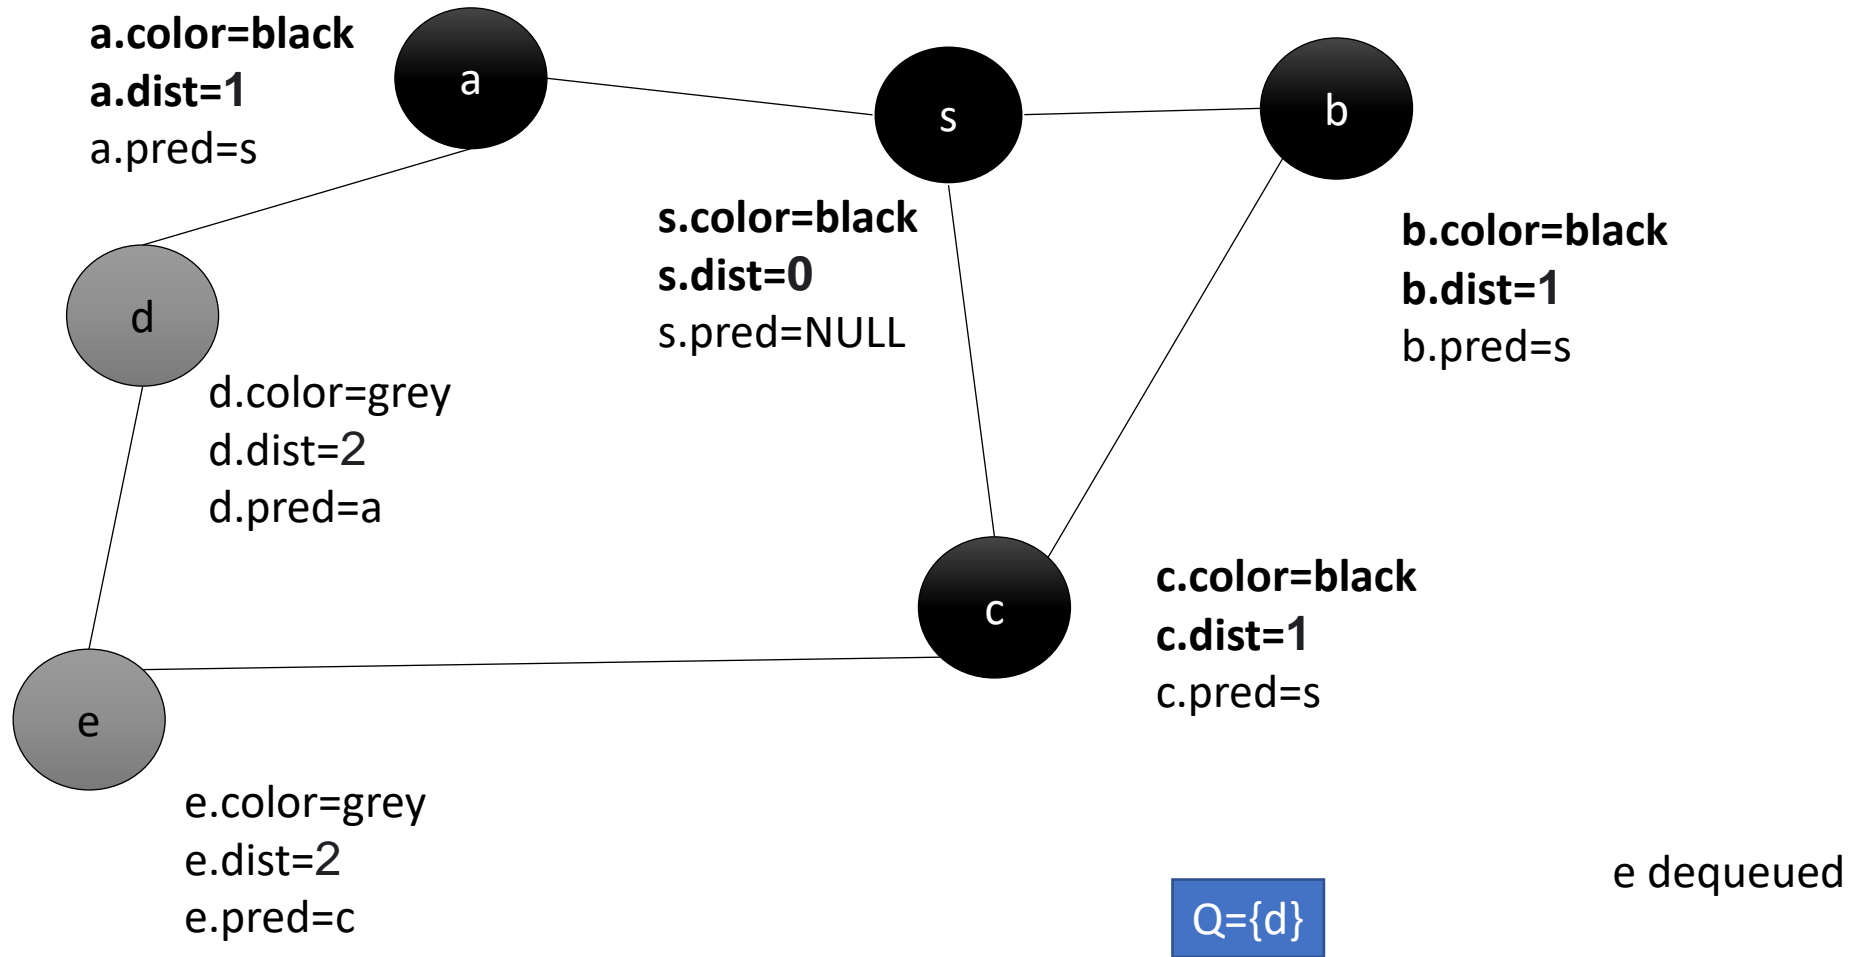
BFS Example



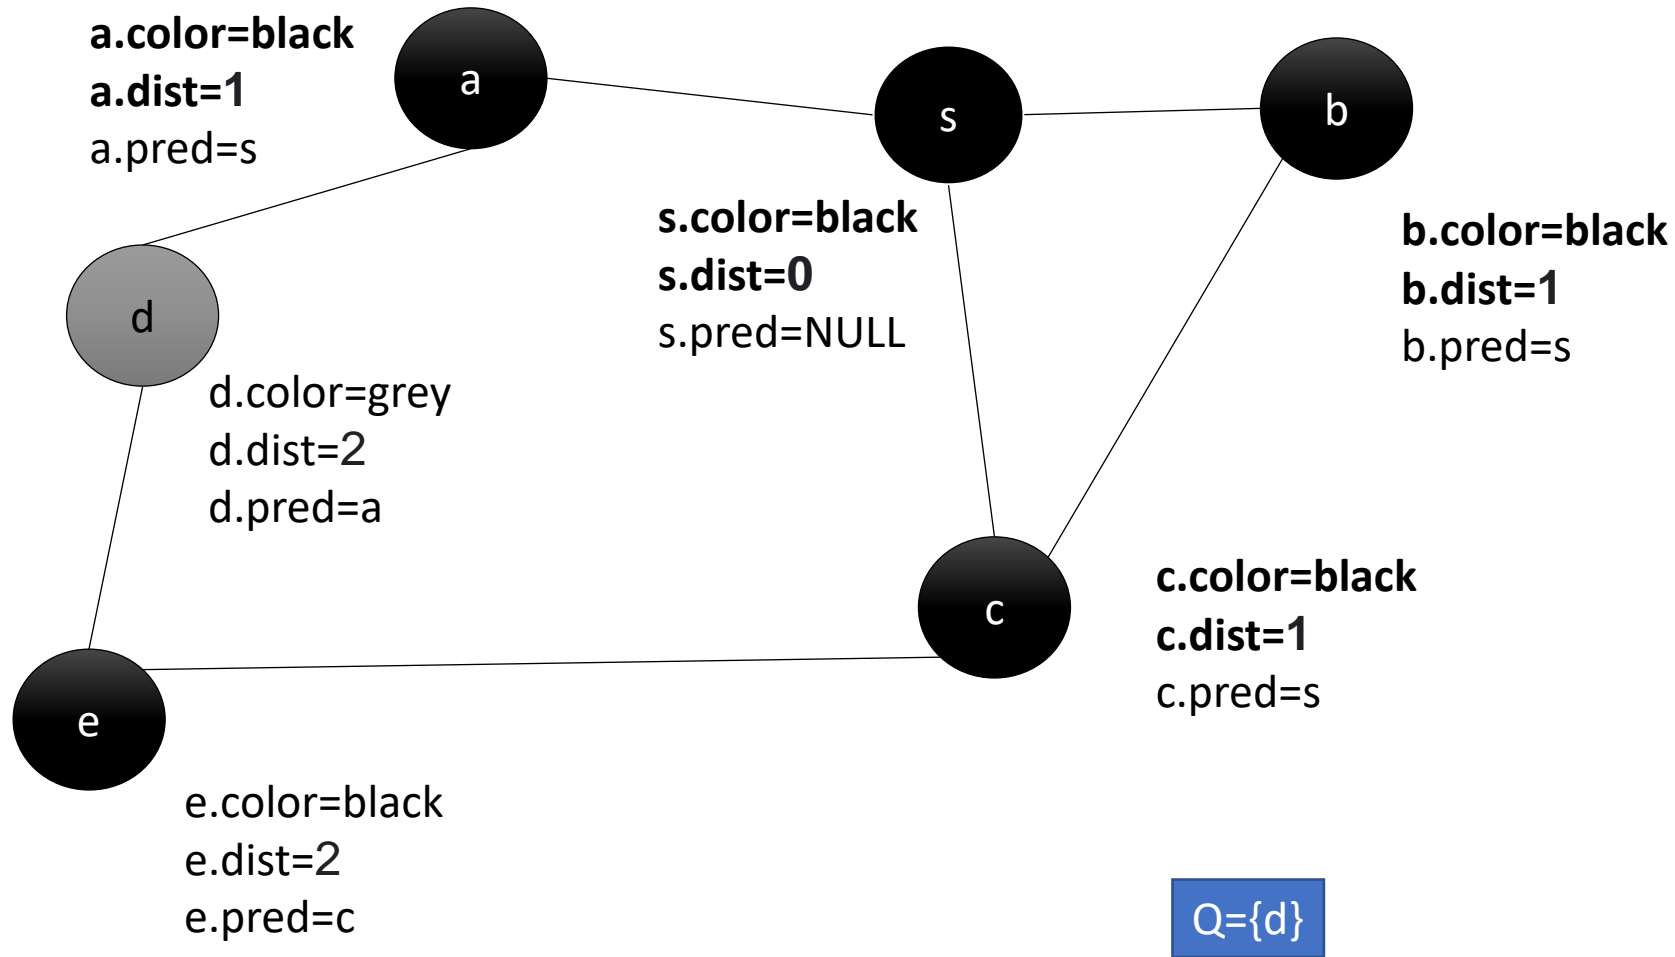
BFS Example



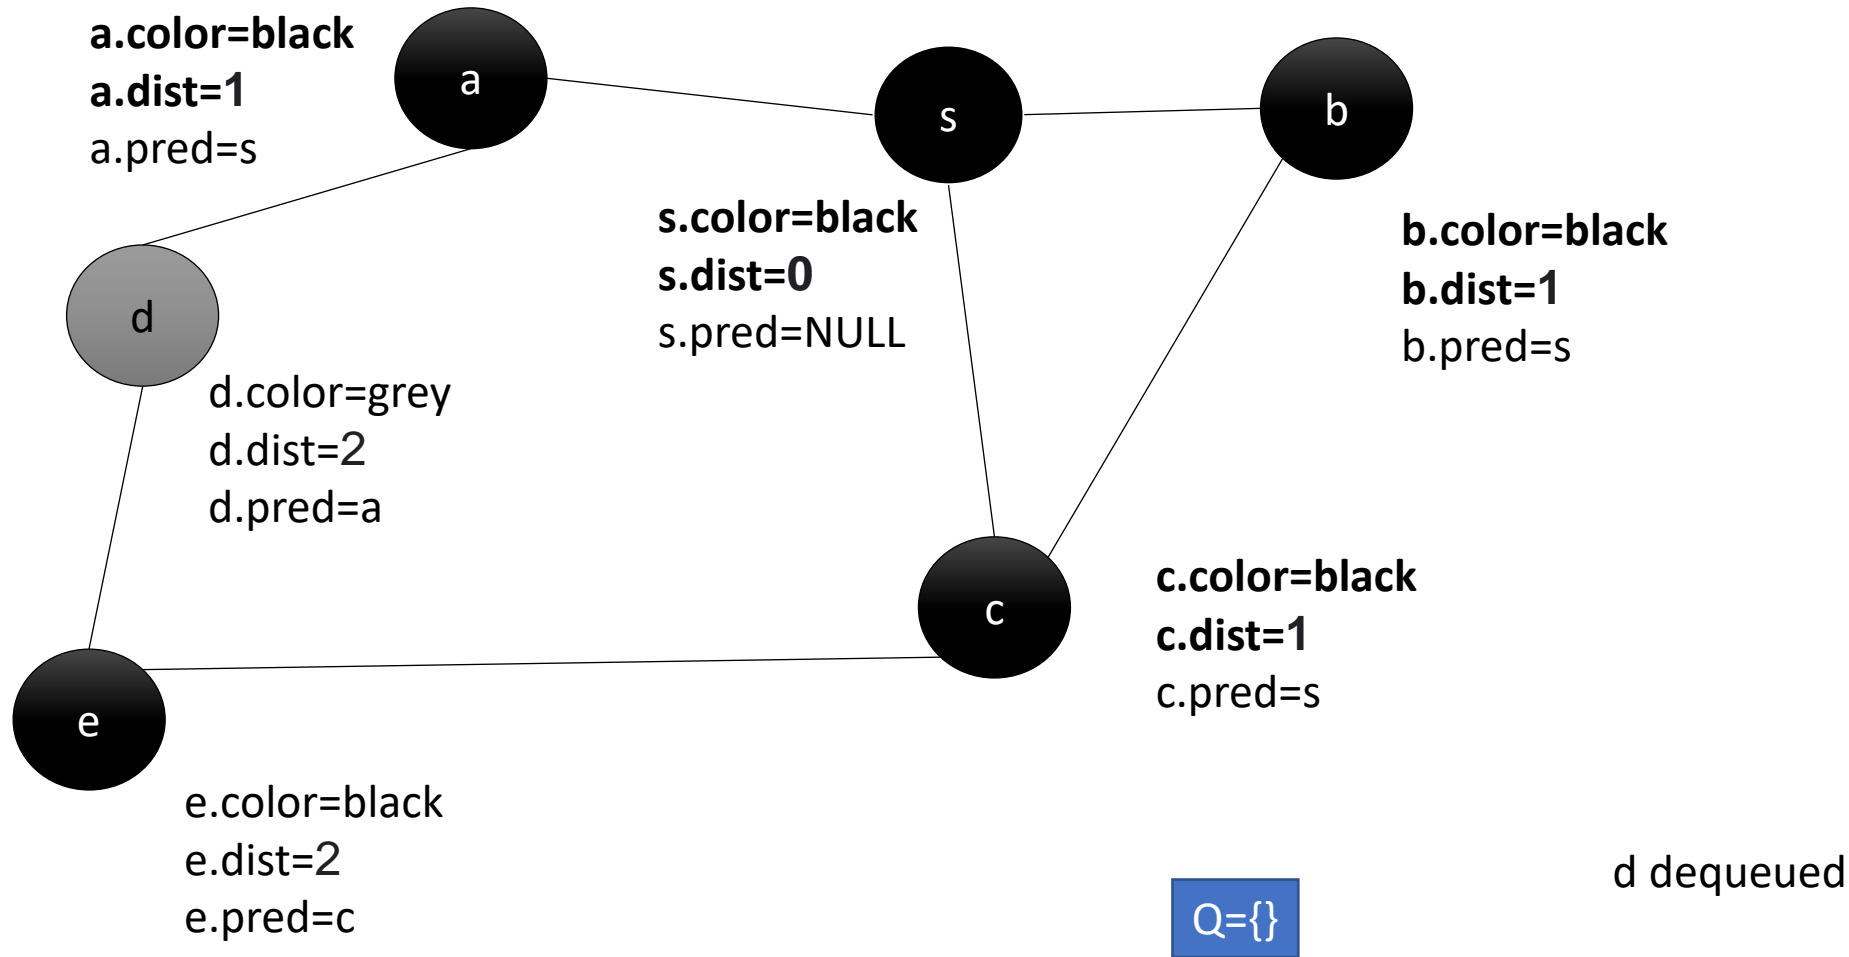
BFS Example



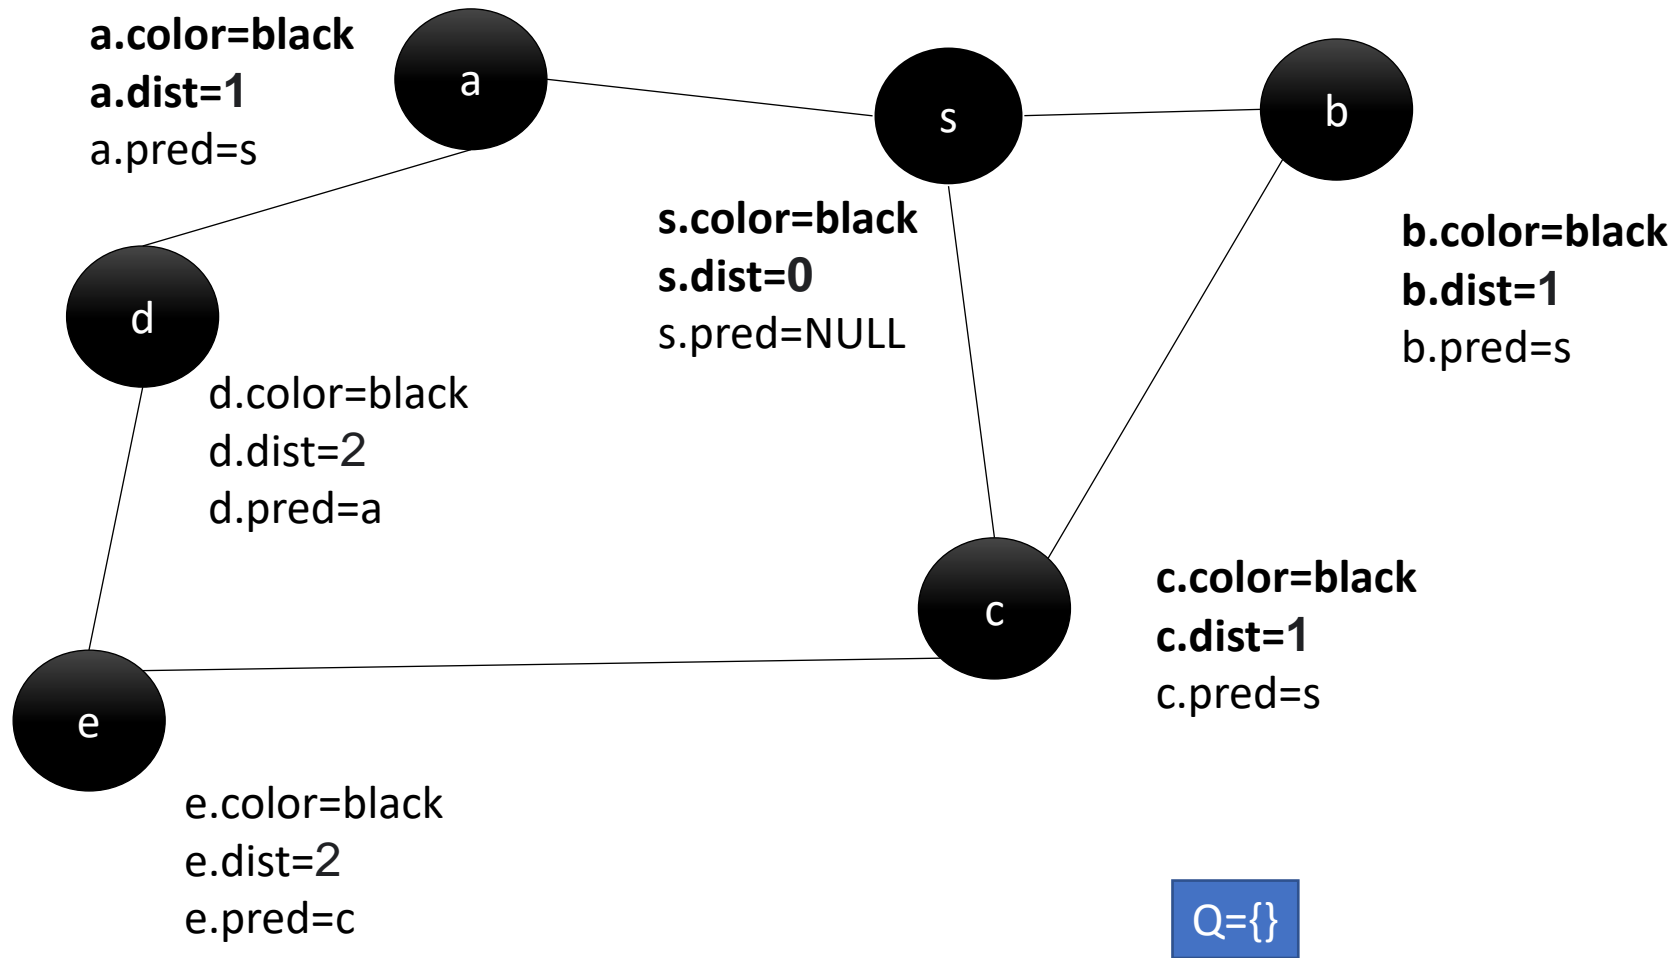
BFS Example



BFS Example



BFS Example



BFS Running time: Exercise

- What is the running time of BFS given a graph represented by an adjacency list?

Applications of BFS

- Find a ***shortest path***, where length is measured in number of edges.
- Find ***connected components*** of an undirected graph
- ***Garbage collection***: traverse the graph of objects reachable from the stack in BFS manner.
 - Enables garbage collector to remove objects that are not reachable.

Depth-first search (DFS)

- Uses notion of *traversal time- starttime and endtime recorded for each vertex*.
- Start at vertex s , mark s as “visited (discovered)”, and label s as current vertex called u
- Travel along an ***arbitrary*** edge (u,v) .
 - If edge (u,v) leads to an already visited vertex v , return to u
 - If vertex v is unvisited, move to v , paint v “visited”, set v as our current vertex, and repeat the previous steps

Depth-first search (DFS)

- Eventually, you get to a point where all edges from u lead to visited vertices
- Backtrack until we get back to a previously visited vertex v
- v becomes our current vertex, and we repeat the previous steps
- Uses a ***stack***.

DFS

- *Uses backtracking*: advance if possible, backtrack if not possible to advance further.
- *Organizes* vertices by *start/end* times.
- Classifies edges as either ***tree*** or ***back edges***.

DFS algorithm

DFS(G):

##Initialize all vertices

For $u \in G.V$ **Do** *#undiscovered*

$u.color = white$

$u.pred = NULL$

$time = 0$ #reset time counter

##Visit all vertices

For $u \in G.V$ **Do**

If $u.color = white$ **Then**

 DFS-VISIT(G,u)

DFS-VISIT(G,u):

##Initialize DFS

$u.color = gray$ *#discovered*

time = time + 1

$u.starttime = time$ #discovery time

##Visit all children-recursively

For $v \in u.adj$ **Do**

If $v.color = white$ **Then**

$v.pred = u$


 DFS-VISIT(G,v)

$u.color = black$ *#finished/processed*


time = time + 1

$u.endtime = time$ #finishing time


DFS Example




DFS Example




DFS Example




DFS Example




DFS Example




DFS Example




DFS Example




DFS Example




DFS Example



DFS Example



DFS Example




DFS: edge classification

- Traversing an **undirected graph** in DFS reveals only ***tree and back edges***.
- Traversing a **directed graph** reveals ***tree, back, forward, and cross edges***.


DFS: edge classification

- **Tree edge**: gray to white
 - Edges in depth first forest.
 - (u,v) is a tree edge if **v is visited for the first time from u .**
 - Form a **tree** containing **each vertex visited in G .**
- **Back edge**: gray to gray
 - Non-tree edge from descendant to ancestor.
 - (u,v) is a back edge if **v appears before u** based on start-time and there is a **path from v to u .**



DFS: edge classification


- **Forward edge:** gray to black
 - Non-tree edge from ancestor to descendant.
 - (u,v) is a forward edge if **v appears after u** based on start-time and there is a **path from u to v**.
- **Cross edge:** gray to black
 - Between trees or subtrees.
 - (u,v) is a cross edge if **v is neither an ancestor nor a descendant of u**.



Most algorithms use **tree and back edges**.

Examples? Later.

Exercise: Edge Classification



Identify each of the following, if any:

- a) Tree edges
- b) Back edges
- c) Forward edges
- d) Cross edges

DFS Applications

- Find **connected components** (in an undirected graph)
- **Check digraph** for cycles.
 - A graph, G , has a cycle it has at **least one back edge**.
- **Graph mining and social network analysis** e.g. relationship discovery, categorization, recommendation(e.g. people you may know, videos you may like), etc.
 - Example: [link](#)
 - Chakrabarti, D. (2011). Graph Mining. In: Sammut, C., Webb, G.I. (eds) Encyclopedia of Machine Learning. Springer, Boston, MA. https://doi.org/10.1007/978-0-387-30164-8_350
- **Topological sorting**: Find a linear order that refines a given partial order (e.g. organizing courses by prerequisites, grooming and dressing up, implementing a recipe)

BFS vs. DFS

- BFS
 - BFS visits all vertices that are reachable from the start node s and ***returns one search tree***.
- DFS
 - Visits all vertices in the graph and ***may return multiple search trees***.
 - DFS classifies edges as ***tree, forward, back, and cross edges***.

Summary

- Definitions: graphs, vertices, edges, cycles, subgraphs, etc.
- Two graph representations: ***adjacency list and adjacency matrix***
- Graph traversal: ***BFS and DFS***
- Applications of graph search algorithms

Next

- DAGs and topological sorting,
- Min-Spanning tree: Prims/Kruskals,
- Shortest Path algorithms: Dijkstras, Floyds